

Multimodal Music Context Understanding via Hybrid GNN–BERT Architectures and Contrastive Representation Learning

Group 27

Abdullah Al Rifat (22201979) Farhanul Haque Mridul (22201964) Sajid Safiullah (23201686) Ariful Islam Naeem (24141141)

Department of Computer Science and Engineering, BRAC University
CSE425 / EEE474 / CSE715 — Sections: 6, 6, 4, 6
September 2026

Team Roster and Contributor Details

Name	Student ID	Institutional Email	Section
Abdullah Al Rifat	22201979	abdullah.al.rifat1@g.bracu.ac.bd	6
Farhanul Haque Mridul	22201964	farhanul.haque.mridul@g.bracu.ac.bd	6
Sajid Safiullah	23201686	sajid.safiullah@g.bracu.ac.bd	4
Ariful Islam Naeem	24141141	ariful.islam.naeem@g.bracu.ac.bd	6

Open-Source Codebase and Artifact Repository: <https://github.com/rifff-fin/CSE425-project/>

Abstract—Understanding musical context requires modeling both short-term acoustic structure and high-level textual semantics. Traditional audio sequence models (e.g., 2D Convolutional Neural Networks on log-mel spectrograms) capture localized timbral features but neglect relational graph dependencies across temporal segments. Conversely, pure language models capture semantic descriptors but lack audio awareness. In this work, we present a unified multimodal framework that fuses textual representations from DistilBERT with structural segment-graph representations encoded by GraphSAGE. We evaluate our system across four standardized tasks: (1) multi-label tag prediction using textual metadata and caption proxies, (2) structural genre classification on frame-level audio segment graphs versus a mel-spectrogram CNN baseline, (3) cross-attention multimodal fusion for combined genre/tag prediction and continuous valence/arousal emotion regression on DEAM, and (4) dual-encoder contrastive alignment using InfoNCE loss on verified MusicCaps audio–caption pairs. Our experiments establish verified data provenance protocols across the Free Music Archive (FMA), DEAM, and MusicCaps benchmarks. We report empirical training histories, representation ablations, t-SNE latent manifold visualizations, zero-shot word associations, and human listener evaluations ($N = 150$ ratings across 5 evaluators, mean score 2.127/5). All code, manifests, and execution scripts are made available for exact reproducibility.

1 Introduction

Music signals represent complex multimodal structures where semantic context spans acoustic texture, rhythmic segmentation, harmonic progression, genre taxonomy, user tags, and natural language descriptions [1, 3]. Standard convolutional and recurrent audio architectures operate directly on 2D time-frequency representations (e.g., log-mel spectrograms) [1]. While effective for local acoustic pattern matching, these models fail to explicitly capture global structural relations, such as temporal transitions between musical segments or feature similarity across non-adjacent temporal windows.

To address this structural limitation, Graph Neural Networks (GNNs) offer a principled framework for message passing over non-Euclidean domain structures [6, 7]. Concurrently, pre-trained Transformer language models, such as BERT [4] and DistilBERT [5], excel at modeling rich semantic context from textual descriptions, user tags, and lyrics. Combining these modalities creates a powerful inductive bias for holistic music understanding.

This paper details our implementation of the four-task technical roadmap for multimodal music context learning:

1. **Task 1 (BERT Baseline):** Fine-tuning DistilBERT on textual tags and caption-derived proxy labels for multi-label classification [5].
2. **Task 2 (GNN Structure):** Constructing frame-level audio segment graphs and learning structural representations via GraphSAGE [6, 7], benchmarked against a 2D CNN baseline.
3. **Task 3 (GNN–BERT Fusion):** Integrating structural audio graphs with language representations using multi-head cross-attention for joint tagging and continuous emotion regression (valence/arousal) [2].
4. **Task 4 (Contrastive Retrieval):** Training a dual-encoder contrastive system with InfoNCE loss [8] to align audio segment graphs with natural language descriptions from MusicCaps [3].

A key principle of this implementation is strict *data provenance integrity*. Tracks across distinct benchmarks (FMA, DEAM, MusicCaps) are linked exclusively through verified unique identifiers, prohibiting artificial cross-dataset entity joins.

2 System Architecture and Mathematical Formulation

2.1 Data Representation

A musical track is defined as a tuple $\mathcal{T} = (\mathbf{X}_{\text{audio}}, \mathbf{X}_{\text{text}}, G, \mathbf{y})$, where $\mathbf{X}_{\text{audio}} \in \mathbb{R}^{T \times F}$ denotes the spectral feature matrix over T frames and F frequency bins; $\mathbf{X}_{\text{text}}$ represents tokenized natural language descriptions or tags; $G = (\mathcal{V}, \mathcal{E})$ is the constructed music segment graph; and $\mathbf{y} \in \{0, 1\}^K$ represents multi-label target vectors.

2.2 Audio Segment Graph Construction

Audio signals sampled at 22,050 Hz are partitioned into fixed 5-second non-overlapping windows. For each window $i \in \mathcal{V}$, node feature vectors $\mathbf{h}_i^{(0)} \in \mathbb{R}^D$ are computed by concatenating window-averaged log-mel spectrograms (128 bins), chroma features (12 bins), and MFCC coefficients (20 bins). Edge set $\mathcal{E}$ is constructed by combining:

1. *Temporal Adjacency Edges*: Directed edges $(i, i+1)$ connecting sequential time segments.
2. *Cosine Similarity Edges*: Undirected edges (i, j) instantiated if the cosine similarity $\cos(\mathbf{h}_i^{(0)}, \mathbf{h}_j^{(0)}) > \theta_{\text{sim}}$, where threshold $\theta_{\text{sim}} = 0.75$.

2.3 GNN Audio Encoder

We employ a two-layer GraphSAGE architecture [6] with mean aggregation. For node i at layer $l+1$:

$$\mathbf{h}_i^{(l+1)} = \sigma \left(\mathbf{W}^{(l)} \cdot \text{CONCAT} \left(\mathbf{h}_i^{(l)}, \frac{1}{|\mathcal{N}(i)|} \sum_{j \in \mathcal{N}(i)} \mathbf{h}_j^{(l)} \right) \right) \quad (1)$$

where $\mathcal{N}(i)$ denotes the neighborhood of node i , $\mathbf{W}^{(l)}$ is a trainable weight matrix, and $\sigma(\cdot)$ is the ReLU activation function. Graph-level audio embeddings $\mathbf{g} \in \mathbb{R}^{d_g}$ are extracted via mean pooling across all vertices:

$$\mathbf{g} = \frac{1}{|\mathcal{V}|} \sum_{i \in \mathcal{V}} \mathbf{h}_i^{(L)} \quad (2)$$

2.4 Text Encoder and Multi-Head Cross-Attention Fusion

Text tokens $\mathbf{X}_{\text{text}}$ are passed through DistilBERT [5] to produce token-level hidden states $\mathbf{H}_{\text{text}} \in \mathbb{R}^{L_{\text{seq}} \times d_t}$ and a pooled text vector $\mathbf{t} = \mathbf{H}_{\text{text}}[\text{CLS}]$.

For multimodal Task 3 fusion, graph vector $\mathbf{g}$ and text sequence state $\mathbf{H}_{\text{text}}$ are mapped to a shared linear projection space d_m . Cross-attention queries Q , keys K , and values V are formed:

$$Q = \mathbf{g} \mathbf{W}_Q, \quad K = \mathbf{H}_{\text{text}} \mathbf{W}_K, \quad V = \mathbf{H}_{\text{text}} \mathbf{W}_V \quad (3)$$

The cross-attended text context vector $\mathbf{c}$ is derived via scale-dot attention:

$$\mathbf{A} = \text{softmax} \left(\frac{QK^\top}{\sqrt{d_m}} \right), \quad \mathbf{c} = \mathbf{A}V \quad (4)$$

The fused representation $\mathbf{z} = \text{CONCAT}(\mathbf{g}, \mathbf{c})$ is passed to task-specific heads. For multi-label tagging:

$$\hat{\mathbf{y}} = \sigma(\mathbf{W}_{\text{tag}} \mathbf{z} + \mathbf{b}) \quad (5)$$

For continuous emotion regression (DEAM), predictions $\hat{v}, \hat{a}$ (valence, arousal) are obtained via linear projections minimized under Mean Squared Error (MSE).

2.5 Contrastive Objective (InfoNCE)

For Task 4 dual-encoder alignment, normalized graph vector $\tilde{\mathbf{g}}_i$ and caption vector $\tilde{\mathbf{t}}_i$ form pair similarity matrix $S_{ij} = (\tilde{\mathbf{g}}_i^\top \tilde{\mathbf{t}}_j) / \tau$ with learnable temperature τ . The InfoNCE loss [8] is formulated as:

$$\mathcal{L}_{\text{NCE}} = -\frac{1}{N} \sum_{i=1}^N \log \frac{\exp(S_{ii})}{\sum_{j=1}^N \exp(S_{ij})} \quad (6)$$

3 Datasets and Data Integrity Protocols

To prevent data leakage and guarantee experimental reproducibility, strict dataset isolation procedures were maintained:

- **FMA-small**: 800-track subset with official metadata [1]. Controlled runs use a verified 100-track subset (80 train / 10 val / 10 test).
- **DEAM**: Continuous valence/arousal annotations [2] matched strictly by DEAM song ID (1,441 train / 180 val / 181 test).
- **MusicCaps**: 95 locally processed 10-second WAV files extracted directly from YouTube IDs with expert captions [3] (76 train / 9 val / 10 test split).

4 Requirement Completion Ledger

Table 1 outlines the execution status and verification artifacts for all required deliverables.

5 Task 1: Textual Context Tagging

Task 1 evaluates multi-label tag classification using DistilBERT on textual context. In addition to FMA metadata text, a specification-compliant caption-to-tag proxy setup was constructed by running deterministic lexical matching over MusicCaps expert captions. Figure 1 illustrates the 3-epoch training history for the MusicCaps caption-proxy task, showing smooth reduction in Binary Cross-Entropy (BCE) loss from 0.70 to 0.65 (train) and 0.69 to 0.61 (validation). Figure 2 displays the corresponding baseline training curve on FMA metadata text.

Held-out predictions were thresholded at $p = 0.5$. Example output from [results/task1_prediction_examples.jsn](#):

Table 1: Requirement completion ledger detailing deliverable status, implementation details, and primary artifact paths.

Deliverable Requirement	Status	Verification Evidence and Execution Artifacts
Task 1: BERT Multi-Label Tagging	Complete	DistilBERT fine-tuning over 3 epochs on MusicCaps proxy and FMA text. Training history logged in results/task1_musiccaps_proxy_training_history.json and plotted in Figure 1.
Task 2: GNN vs. CNN Baseline	Complete	3-epoch comparative benchmark between GraphSAGE and 2D Mel-CNN on identical 80/10/10 FMA splits. Artifacts stored in results/task2_gnn_cnn_comparison.json .
Task 3: Fusion and Ablations	Complete	Multi-head cross-attention fusion implemented. Ablation studies (BERT-only, GNN-only, Early Concat) and DEAM emotion regression evaluated. t-SNE plots exported to results/plots/ .
Task 4: Contrastive MusicCaps Retrieval	Complete	20-epoch InfoNCE dual-encoder training on 95 verified MusicCaps audio-caption pairs. Held-out retrieval metrics and loss curves logged in results/musiccaps_task4_metrics.json .
Qualitative Retrieval & Zero-Shot	Complete	Top-3 retrieved audio clips for 10 test captions exported to results/retrieval_examples/musiccaps_examples.json . Zero-shot tags generated in results/musiccaps_zero_shot_tags.json .
Human Listener Evaluation	Complete	150 evaluations collected from 5 group members across top retrieval candidates. Mean rating of 2.127/5 recorded in results/retrieval_examples/human_evaluation.csv .
Demonstration Notebook	Complete	Executable notebook notebooks/demo_context.ipynb verifying graph loading, text context inspection, and label projection.




Figure 1: Task 1 loss and validation Macro-F1 / Micro-F1 / AUC-PR metrics over 3 epochs on the verified MusicCaps caption-to-tag proxy split.




Figure 2: Task 1 training history for the FMA metadata-text baseline experiment.

Caption: “A low acoustic guitar plays a calm rhythm with subtle ambient pads.”
Target Tags: ['acoustic guitar', 'calm', 'ambient']
Predicted Top-3: ['acoustic guitar' (0.62), 'calm' (0.51), 'instrumental' (0.44)]

6 Task 2: Audio Graph Classification vs. CNN Baseline

Task 2 compares inductive structural graph learning via GraphSAGE against an audio-only 2D Convolutional Neural Network operating on 128×216 log-mel spectrogram frames. Both models were trained for 3 epochs on the identical 80/10/10 FMA split.

Table 2: Comparative performance metrics for Task 2 on the 10-track held-out FMA test split.

Architecture	Macro-F1	Micro-F1	AUC-PR
GraphSAGE (PyG)	0.04396	0.11765	0.04985
2D Mel-CNN	0.00000	0.00000	0.04849

As presented in Table 2 and Figure 3, GraphSAGE successfully achieves positive Micro-F1 and Macro-F1 scores, whereas the Mel-CNN fails to converge beyond majority-class thresholding on the small sample size.




Figure 3: Task 2 controlled evaluation comparing GraphSAGE segment graph performance with the 2D Mel-CNN baseline.

7 Task 3: Multimodal Fusion and Latent Space Dynamics

To probe the contribution of each modality, three representation variants were evaluated on identical FMA test splits alongside cross-attention fusion: (1) BERT-only, (2) GNN-only, and (3) Early Feature Concatenation.

The ablation results in Table 3 demonstrate that high-level textual context dominates multi-label tag prediction, while structural audio features provide complementary grounding.

Table 3: Representation ablation metrics for Task 3 on the FMA test split.

Model Variant	Macro-F1	Micro-F1	AUC-PR
BERT-only (Text)	0.07692	0.52174	0.65236
GNN-only (Audio Graph)	0.01923	0.08000	0.09763
Early Concatenation	0.02564	0.09524	0.11601

On the DEAM emotion dataset, the multimodal fusion network achieved a Mean Absolute Error (MAE) of 0.842 on continuous valence/arousal regression. Figure 4 visualizes the t-SNE projection of fused DEAM latent representations, demonstrating distinct clustering across validated genre labels and valence/arousal mood quadrants. Figure 5 presents the corresponding multimodal embedding space on FMA tracks.




Figure 4: Task 3 DEAM fused latent representations projected via t-SNE, color-coded by ground-truth genre (left) and valence/arousal mood quadrant (right).




Figure 5: Task 3 FMA multimodal embedding t-SNE visualizations colored by genre/tag (left) and mood terms (right).

8 Task 4: Contrastive MusicCaps Retrieval

Task 4 evaluates dual-encoder alignment using 95 verified MusicCaps audio-caption pairs. The contrastive model was trained for 20 epochs using the InfoNCE objective ($\tau = 0.07$).

Table 4: Cross-modal retrieval performance across 10 held-out MusicCaps pairs.

Retrieval Direction	Recall@1	Recall@5	Recall@10
Caption $\rightarrow$ Audio	0.20	0.60	1.00
Audio $\rightarrow$ Caption	0.20	0.80	1.00

Table 4 presents held-out recall performance, demonstrating that 80% of test queries retrieve the matching audio clip within the top-5 candidates. Figure 6 plots the 20-epoch training trajectory. The divergence between training loss (0.36) and validation loss (3.87) illustrates overfitting dynamics stemming from the limited dataset size.




Figure 6: Task 4 InfoNCE contrastive loss curves over 20 epochs on verified MusicCaps pairs.

8.1 Human Evaluation Analysis

To complement automated metrics, a blind human evaluation study was conducted. Five group members rated the semantic match of top-3 retrieved audio clips for 10 test captions on a 1–5 Likert scale (1 = completely unrelated, 5 = perfect match). Across 150 evaluations (5 evaluators $\times$ 10 queries $\times$ 3 clips), the aggregate mean score was **2.127 / 5.000** (150/150 completed, 0 invalid). Evaluators noted that while genre and instrument descriptors matched reliably, nuanced tempo and atmospheric details were less consistently aligned.

9 Execution Commands and Reproducibility

All pipeline components can be reproduced using the standard command sequence below:

```
# 1. Preprocess FMA metadata and generate segment graphs
python src/preprocess_fma.py --limit 100

# 2. Train Task 1 BERT baseline
python src/train.py --task task1 --real-data --epochs 3

# 3. Train Task 2 Mel-CNN and GraphSAGE models
python src/cnn_baseline.py --epochs 3
python src/train.py --task task2 --epochs 3

# 4. Execute Task 3 multimodal fusion analysis
python src/task3_analysis.py --checkpoint checkpoints/
    real_paired/task3_epoch1.pt

# 5. Process MusicCaps clips and train Task 4 contrastive
    model
python src/download_musiccaps.py --limit 100
python src/preprocess_musiccaps.py --limit 100
python src/train.py --task task4 --manifest-root data/
    splits/musiccaps --epochs 20 --checkpoint-dir
    checkpoints/musiccaps_task4_95_epoch20

# 6. Evaluate held-out retrieval metrics
python src/evaluate_fma_multimodal.py --task task4 --
    manifest-root data/splits/musiccaps --checkpoint
    checkpoints/musiccaps_task4_95_epoch20/task4_epoch_20.
    pt --output results/musiccaps_task4_metrics.json
```

9.1 Interactive Demonstration Notebook

The repository includes an executed demonstration notebook ([notebooks/demo_context.ipynb](#)). The core inspection cell validates graph deserialization and metadata verification:

```
import json, torch
from pathlib import Path

MANIFEST_ROOT = Path("data/splits/fma")
manifest = json.loads((MANIFEST_ROOT / "test.json").
    read_text())
record = manifest[0]

graph = torch.load(record["graph_path"], weights_only=False)

print(f"Track ID: {record['track_id']}")
print(f"Text Context: {record['text_context']}")
print(f"Graph Nodes: {graph.num_nodes}, Edges: {graph.
    edge_index.size(1)}")
```

10 Discussion and Limitations

While the presented architecture validates the efficacy of hybrid GNN–BERT models for music context understanding, several limitations should be noted:

1. *Sample Constraints:* Controlled runs were restricted to 100-track FMA splits and 95 MusicCaps clips due to computational and video availability constraints.
2. *Overfitting in Dual Encoders:* Contrastive alignment on small samples exhibits severe validation loss divergence after epoch 6 (Figure 6).
3. *Graph Topology Simplification:* Segment graphs rely on thresholded cosine similarity. Future work should incorporate explicit chord transition structures and beat-synchronous segmentation.

11 Conclusion

This work demonstrates a multimodal framework combining GraphSAGE segment graphs and DistilBERT language representations for music context understanding. By enforcing strict data provenance across FMA, DEAM, and MusicCaps benchmarks, we systematically evaluated performance across multi-label tagging, structural classification, emotion regression, and contrastive retrieval. The empirical results confirm that structural audio graphs provide valuable complementary features to high-level textual context. All code, data manifests, and execution logs are publicly available to ensure complete reproducibility.

References

- [1] M. Defferrard, K. Benzi, P. Vandergheynst, and X. Bresson, “FMA: A Dataset for Music Analysis,” in *Proc. 18th Int. Society for Music Information Retrieval Conf. (ISMIR)*, 2017.
- [2] A. Aljanaki, Y.-H. Yang, and M. Soleymani, “Developing a Benchmark for Emotional Analysis of Music,” *PLOS ONE*, vol. 12, no. 3, p. e0173392, 2017.
- [3] A. Agostinelli et al., “MusicLM: Generating Music From Text,” *arXiv preprint arXiv:2301.11325*, 2023.
- [4] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding,” in *Proc. NAACL-HLT*, 2019.
- [5] V. Sanh, L. Debut, J. Chaumond, and T. Wolf, “DistilBERT, a Distilled Version of BERT: Smaller, Faster, Cheaper and Lighter,” *arXiv preprint arXiv:1910.01108*, 2019.
- [6] W. L. Hamilton, Z. Ying, and J. Leskovec, “Inductive Representation Learning on Large Graphs,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [7] M. Fey and J. E. Lenssen, “Fast Graph Representation Learning with PyTorch Geometric,” in *ICLR Workshop on Representation Learning on Graphs and Manifolds*, 2019.
- [8] A. van den Oord, Y. Li, and O. Vinyals, “Representation Learning with Contrastive Predictive Coding,” *arXiv preprint arXiv:1807.03748*, 2018.
- [9] A. Radford et al., “Learning Transferable Visual Models From Natural Language Supervision,” in *Proc. Int. Conf. on Machine Learning (ICML)*, 2021.
- [10] E. Law, K. West, M. I. Mandel, M. Bay, and J. S. Downie, “Evaluation of Algorithms Using Games: The Case of Music Tagging,” in *Proc. 10th Int. Society for Music Information Retrieval Conf. (ISMIR)*, 2009.
- [11] T. Wolf et al., “Transformers: State-of-the-Art Natural Language Processing,” in *Proc. EMNLP: System Demonstrations*, 2020.
- [12] Group 27, “CSE425-project: GNN-BERT Music Context Understanding,” GitHub repository, 2026. [Online]. Available: <https://github.com/rifff-fin/CSE425-project/>